

Welcome to CS 186, Section 3!

TA: Bryan Munar

OH: Mondays 11-12pm and Thursdays 2:30-3:30pm
(651 Soda)

DISC: Tuesdays 11-12am (136 Barrows) and Wednesdays
10-11am (130 Wheeler)



Announcements and Such

- Project/Homework 2 Released later this week!!
- Sign up to be partners with someone if you would like!

Discussion 3: Join Algorithms!

Overview:

1. Join Algorithms
2. Worksheet exercises
3. More Join Algorithms
4. Worksheet exercises


(A majority of the joins slides are from Michelle! Thank her!)

*Only going through second half of worksheet

```
SELECT * FROM Sailors S, Reserves R  
WHERE S.sid = R.sid;
```

Visualizations

Sailors



Visualizations

Sailors

Page 1

Page 2

Page 3

Page 4

Visualizations

Sailors

Record 1
Record 2
Record 3
Record 4
Record 5

Page 2

Page 3

Page 4

Visualizations

Sailors

Reserves

Record 1

Record 2

Record 3




Record 4

Record 5

Page 2

Page 3

Page 4



Simple Nested Loops Join

Sailors

Reserves

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Simple Nested Loops Join

Sailors

Reserves

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

(name = Bob, sid = 1)

Simple Nested Loops Join

Sailors

Reserves

(name = Bob, sid = 1)

(sid = 3, bid = 6)
(sid = 1, bid = 4)

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Output:

(name = Bob, sid = 1, bid = 4)

Simple Nested Loops Join

Sailors

Reserves

(name = Bob, sid = 1)

(sid = 3, bid = 6)
(sid = 1, bid = 4)
(sid = 1, bid = 7)

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves

(name = Bob, sid = 1)

(sid = 3, bid = 6)
(sid = 1, bid = 4)
(sid = 1, bid = 7)

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

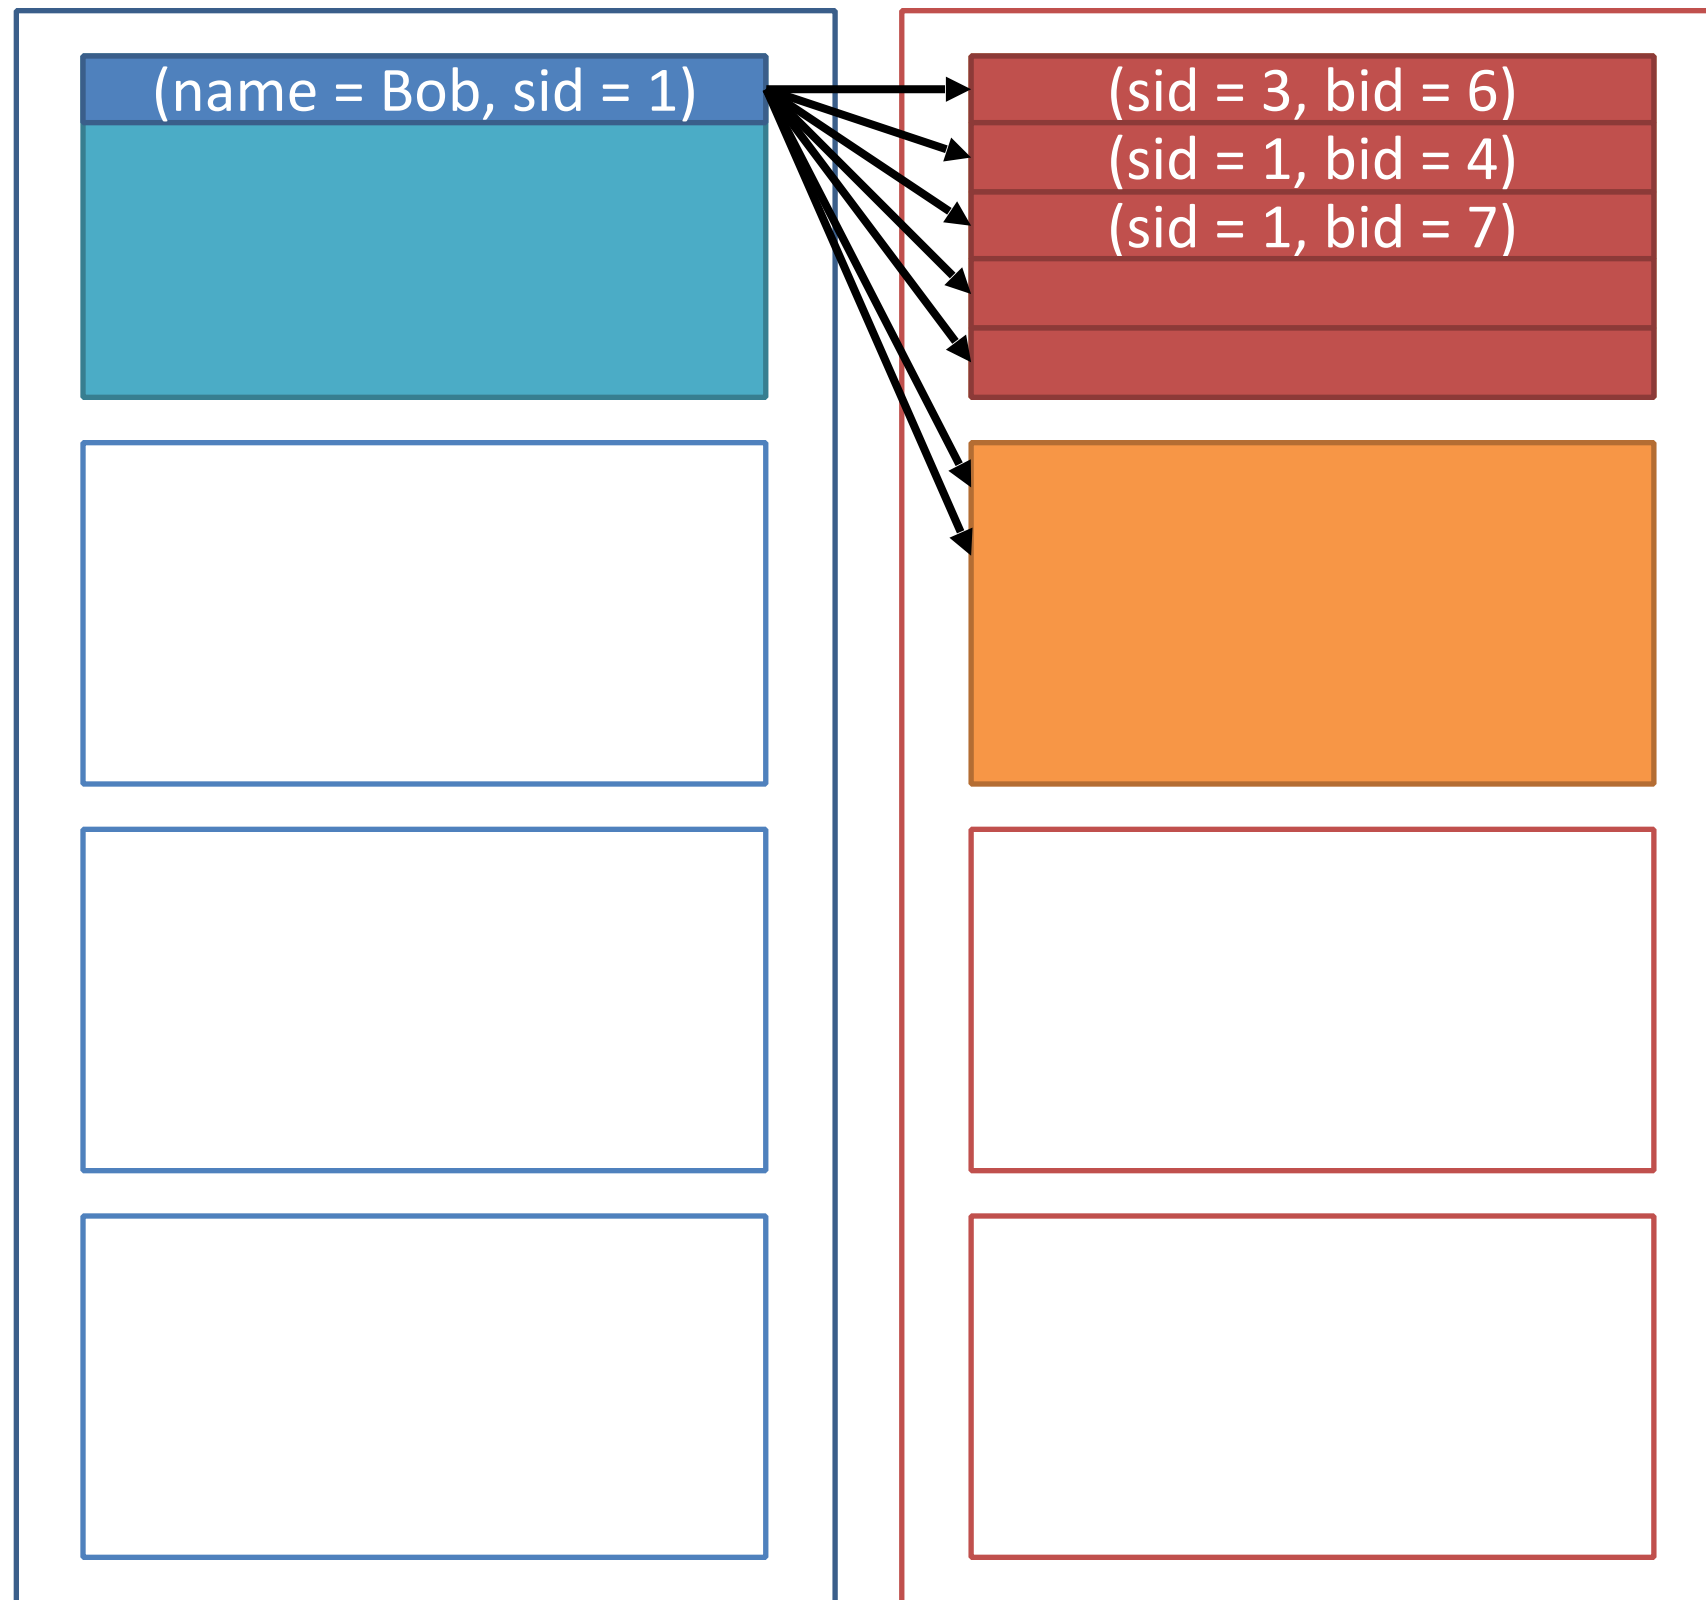
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.


Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

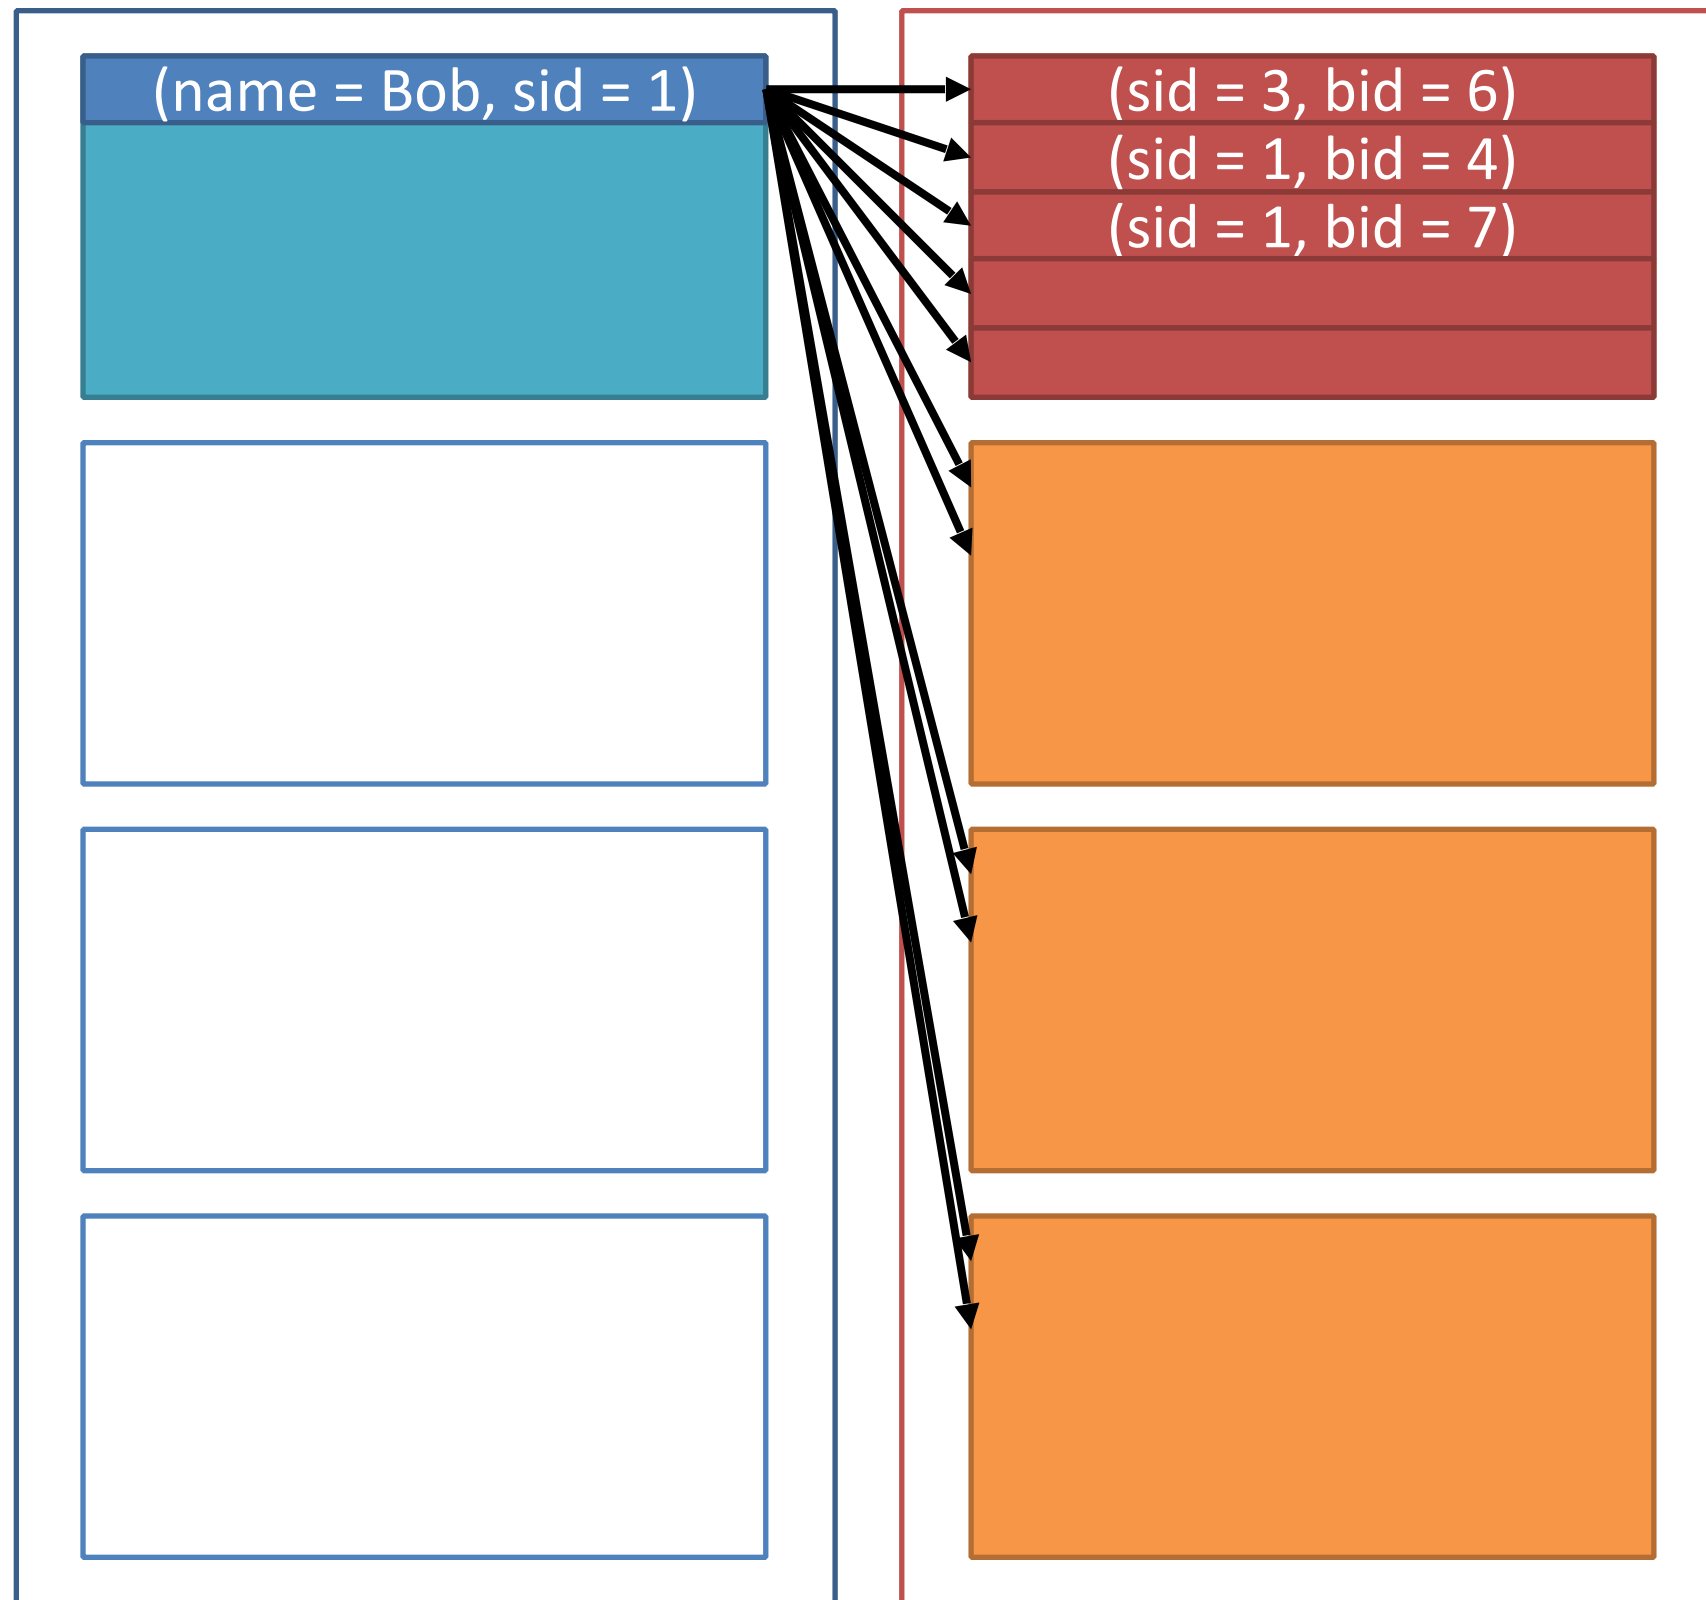
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves

(name = Bob, sid = 1)
(name = Sam, sid = 3)

(sid = 3, bid = 6)
(sid = 1, bid = 4)
(sid = 1, bid = 7)

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Output:


(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Notation: $[S]$ == “# pages in S ” ; $|S|$ == “# tuples in S ”

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.



I/Os:

$$[S] + |S| * [R]$$

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.



Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.



Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

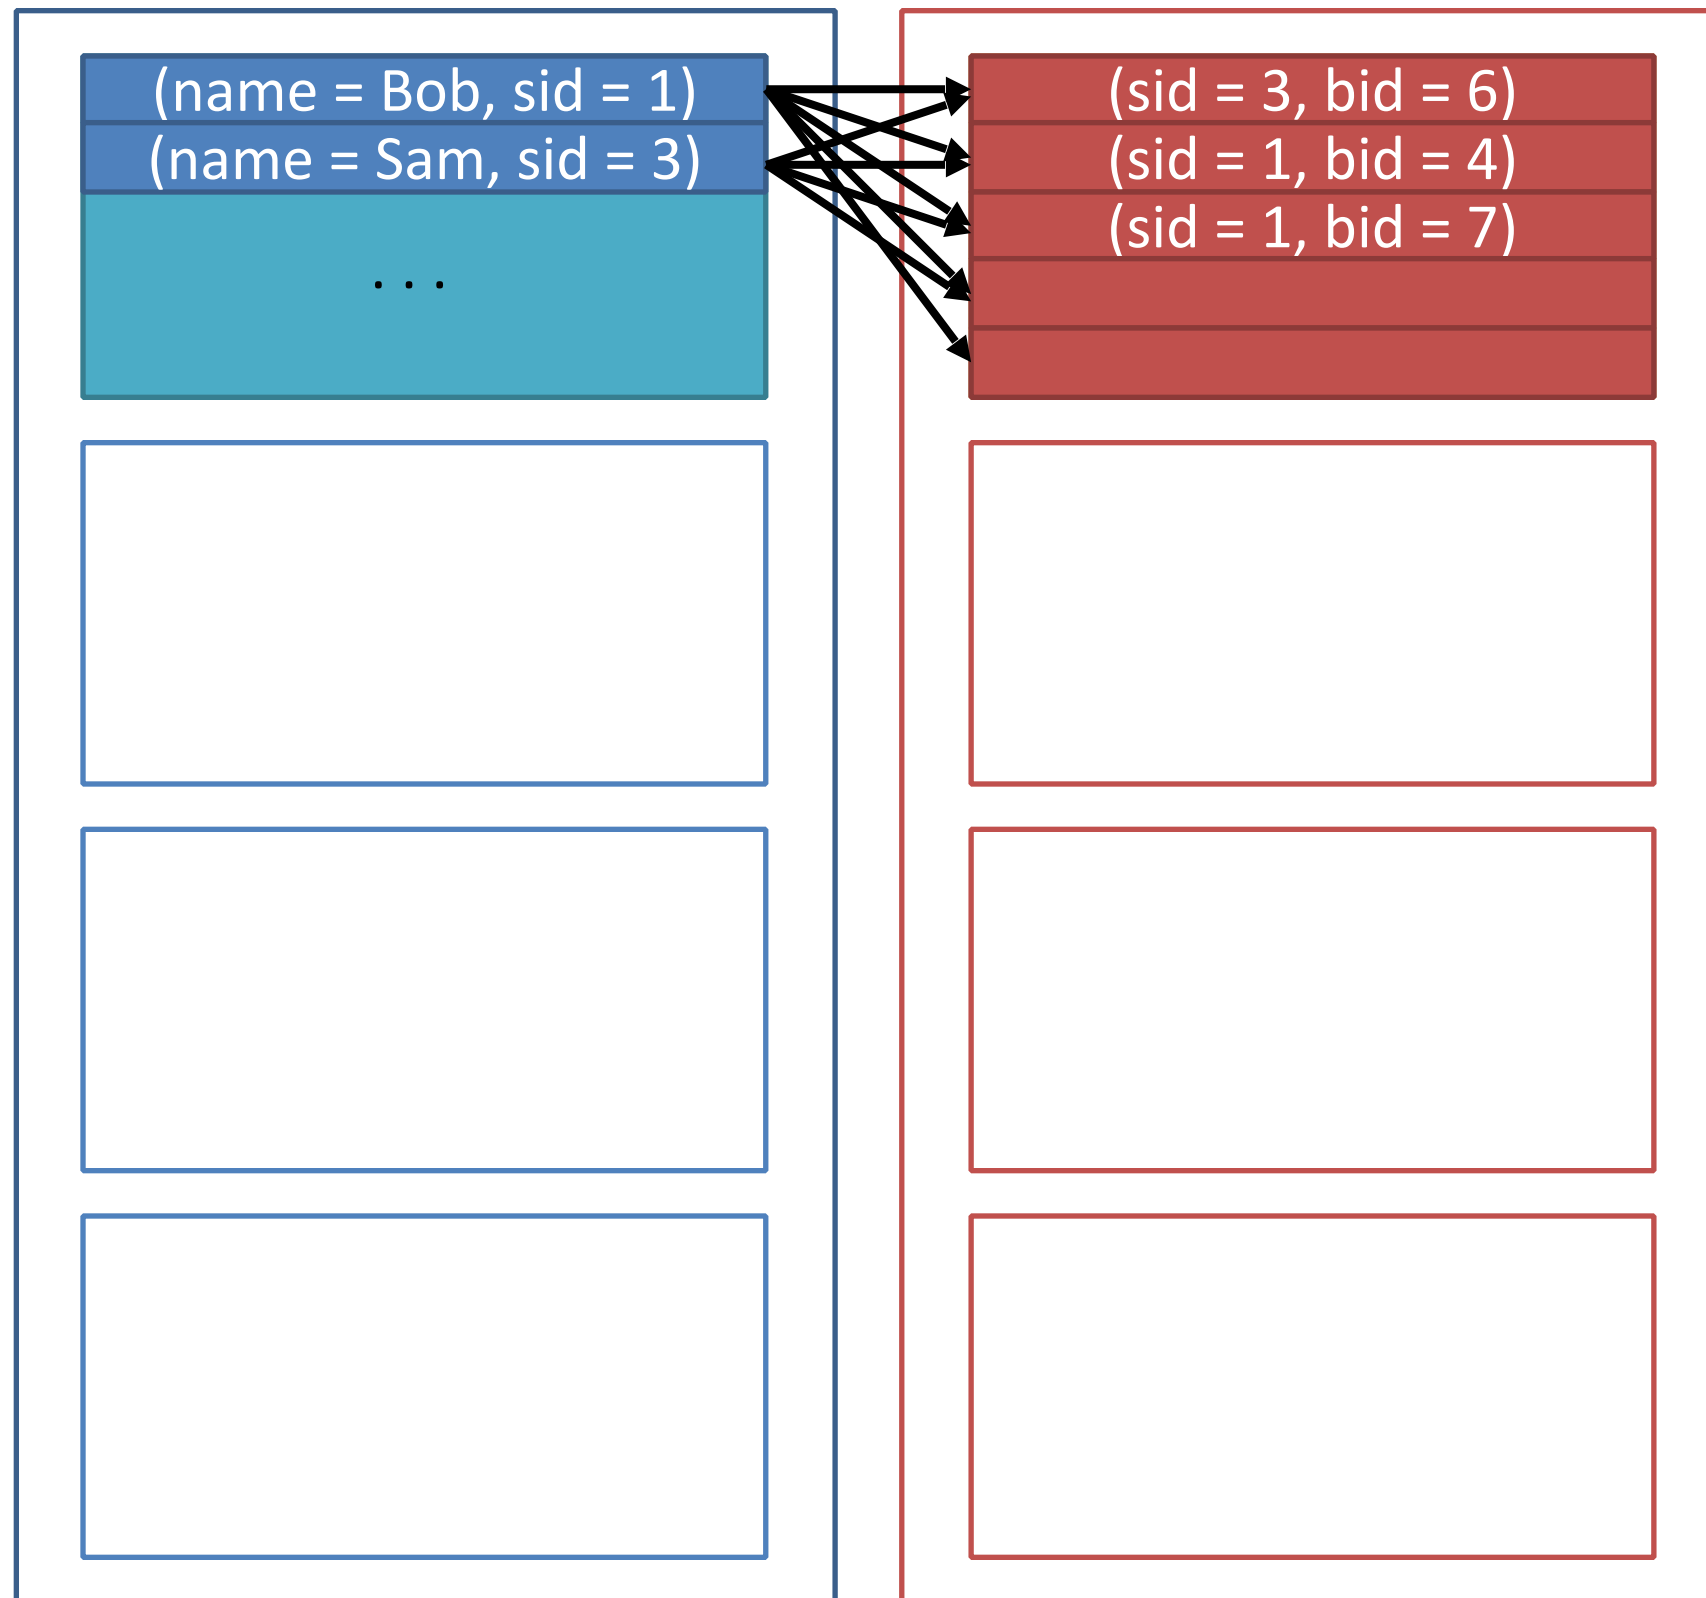
Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.



Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.



Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

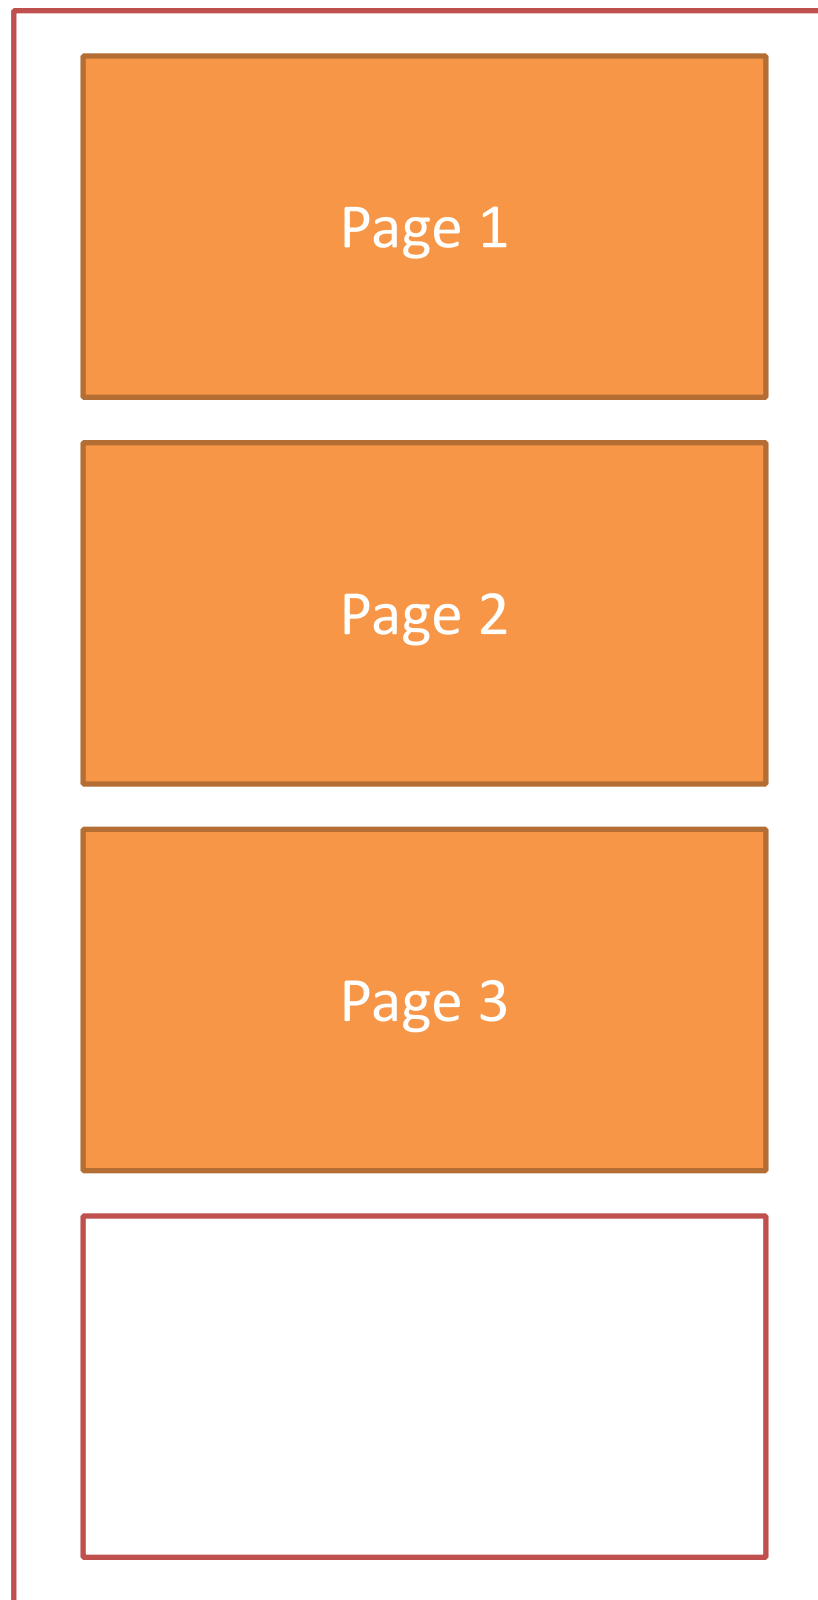

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

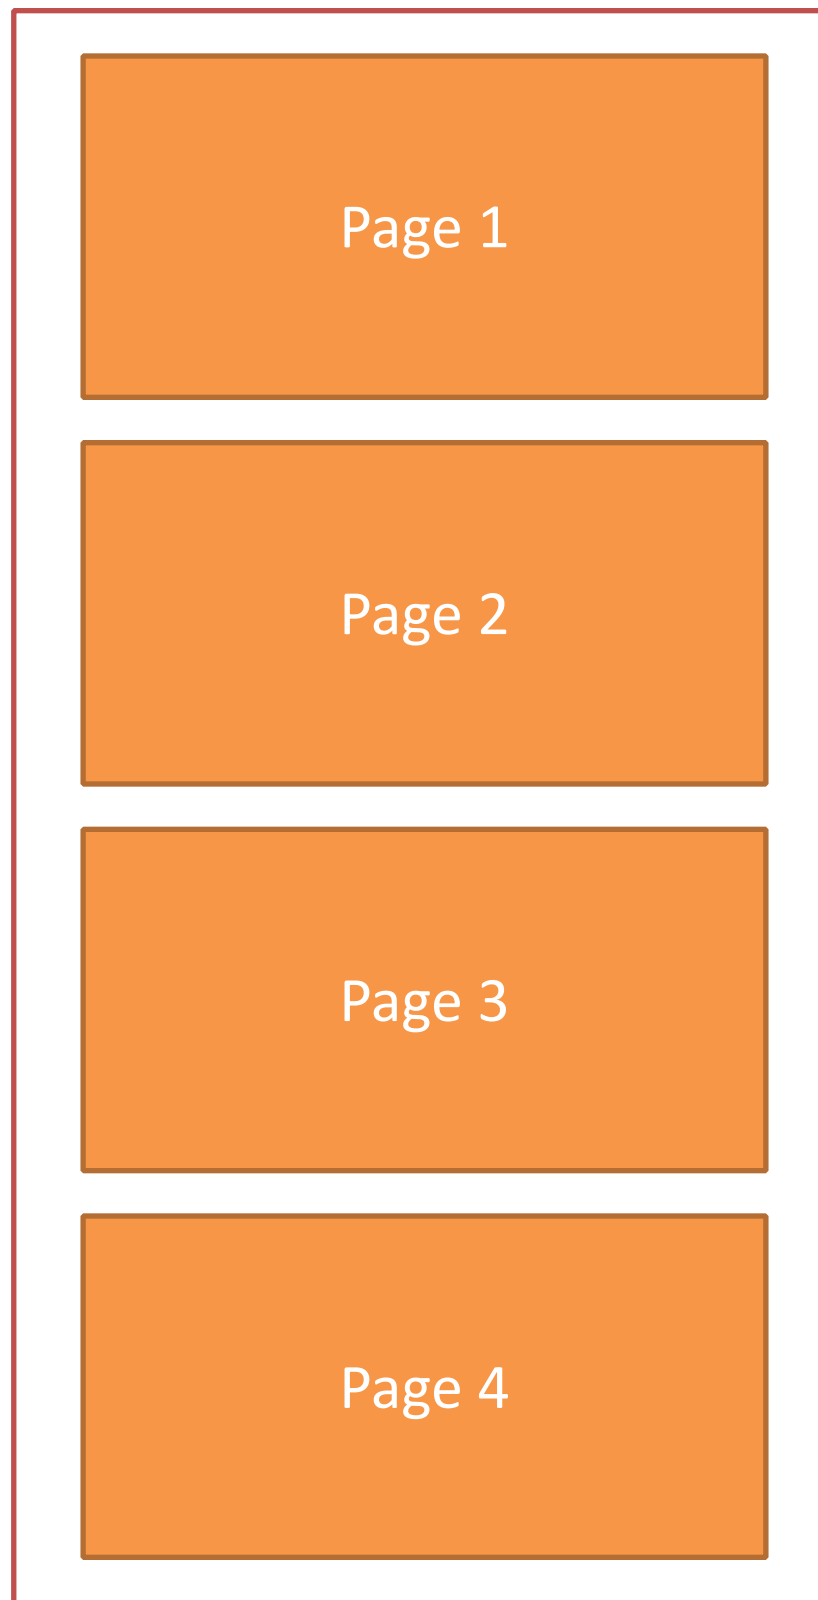

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves

Page 1

Page 2

Key idea:

Take each page of S
and match with each
page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.



Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Output:


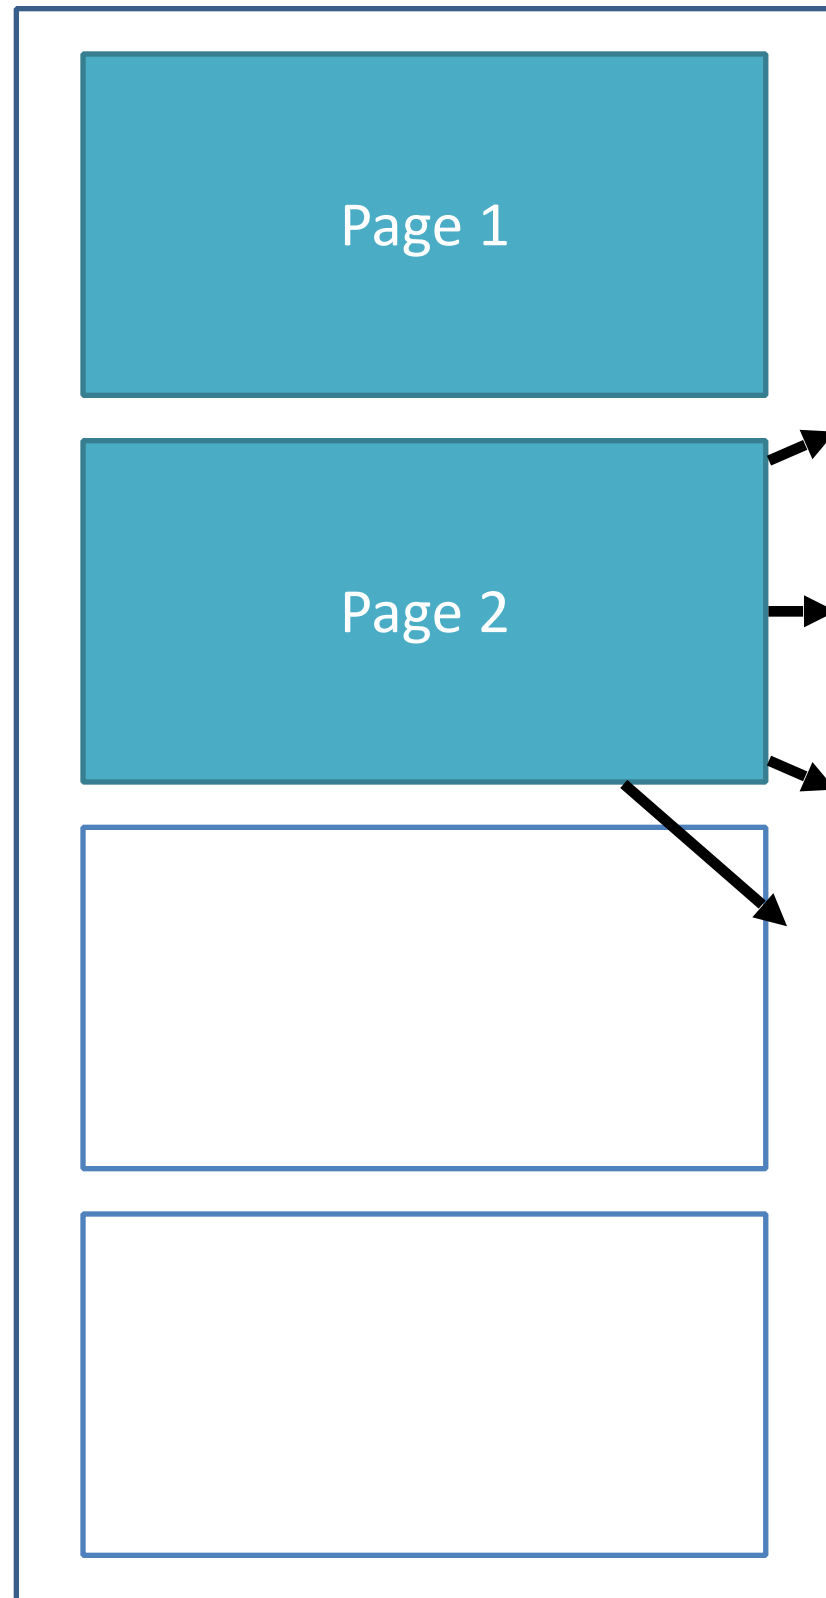
(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Notation: $[S]$ == “# pages in S” ; $|S|$ == “# tuples in S”

Page-Oriented Nested Loops Join

Sailors

Reserves



Do we want the smaller relation as the OUTER or the INNER?

Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.



Compare tuples in each.

$$[S] + [S]*[R]$$

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S
and match with
each page of R.


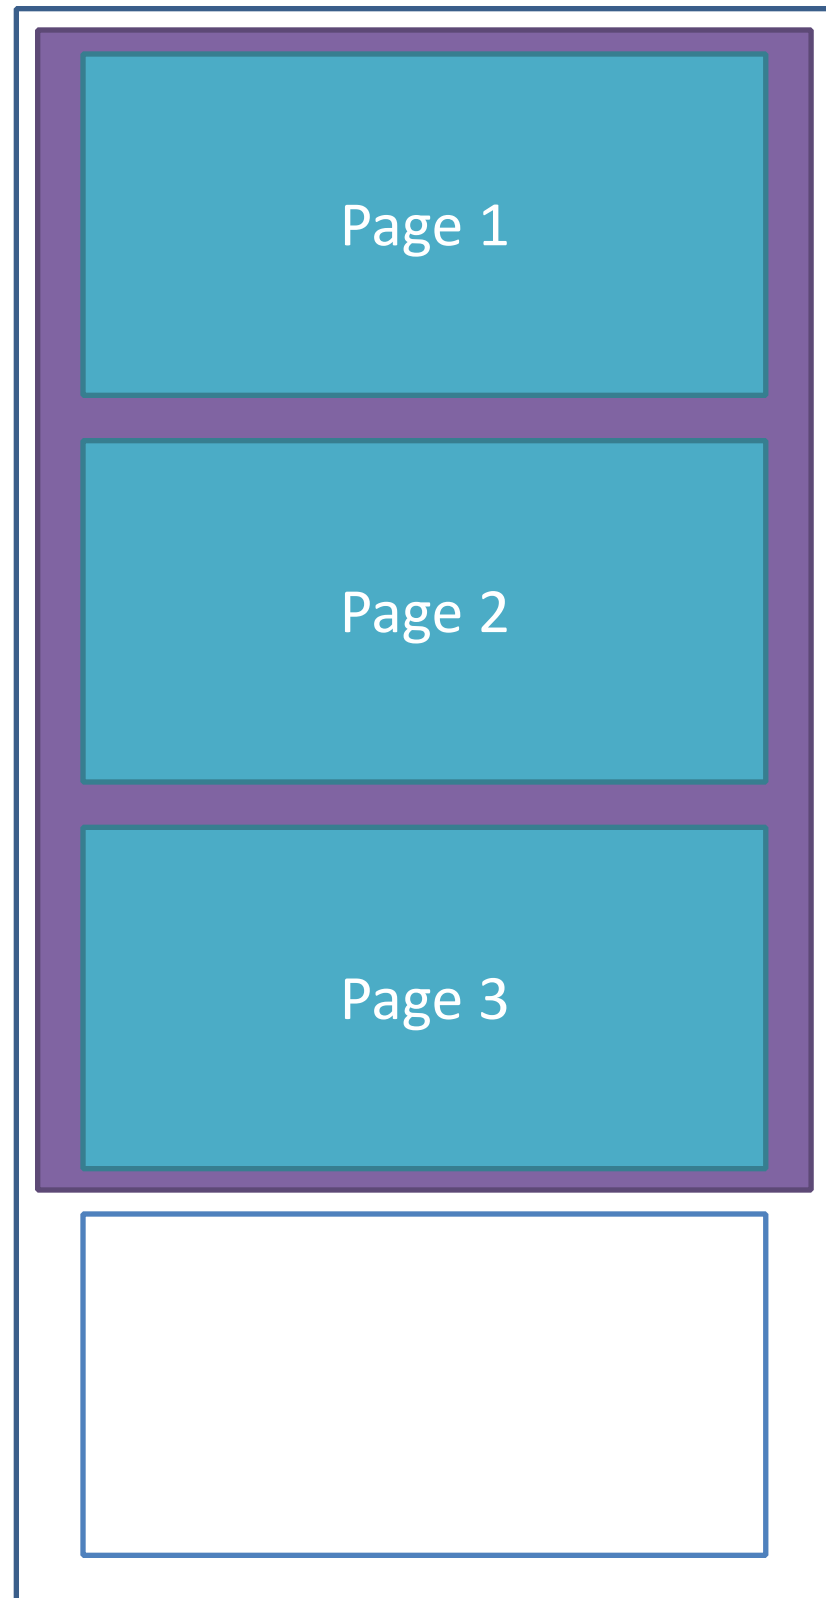
Steps:

1. Get **k** pages of S.
2. Iterate through
each page in R.
3. Compare tuples in
each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S
and match with
each page of R.

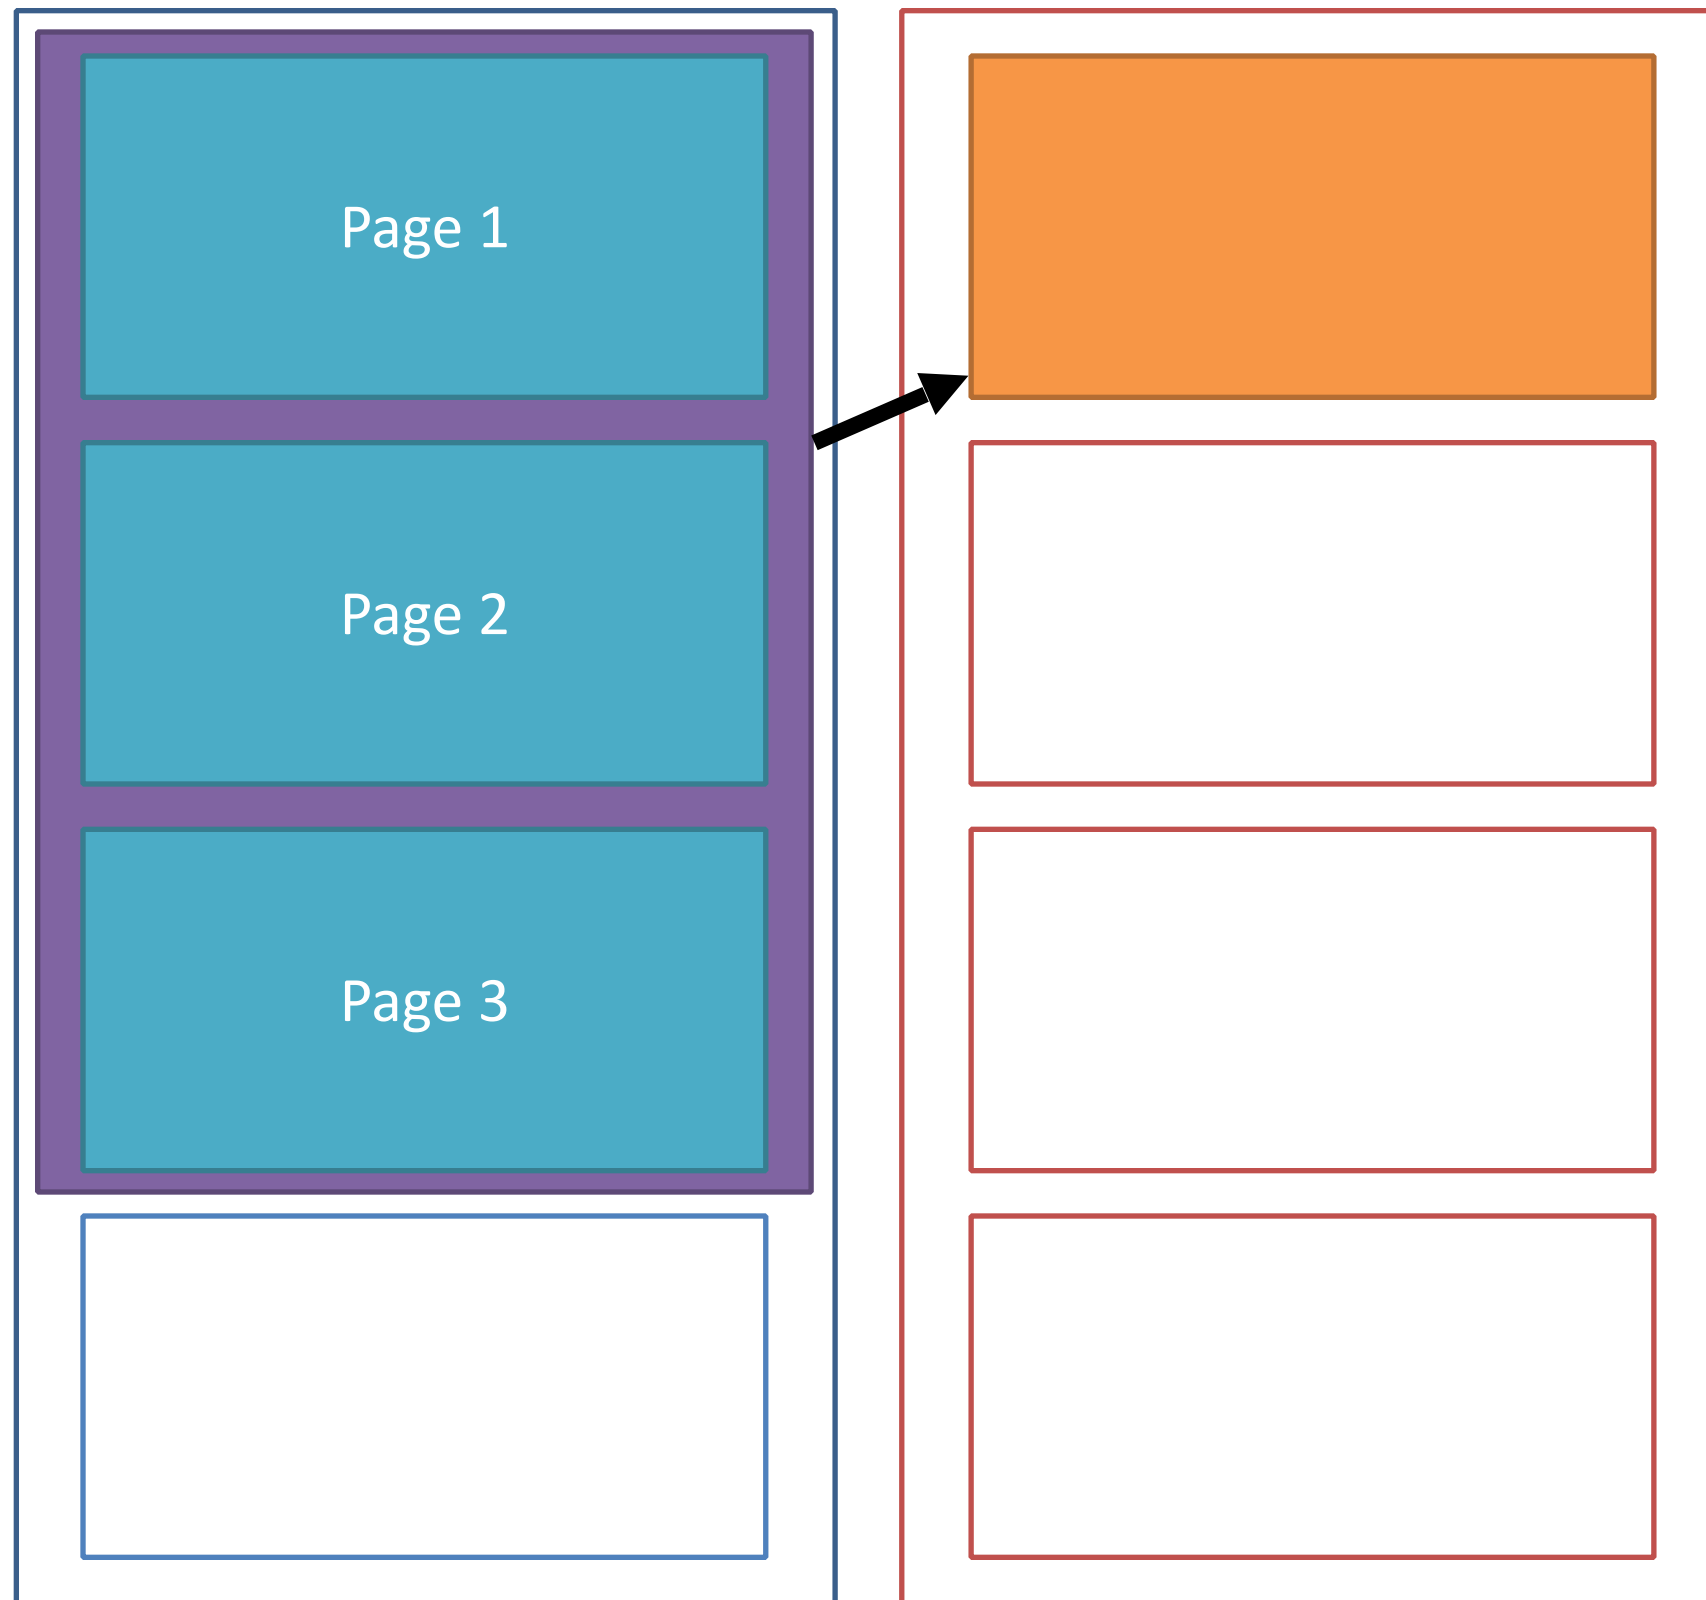
Steps:

1. Get **k** pages of S.
2. Iterate through
each page in R.
3. Compare tuples in
each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S
and match with
each page of R.


Steps:

1. Get **k** pages of S.
2. Iterate through
each page in R.
3. Compare tuples in
each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.


Steps:

1. Get **k** pages of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.

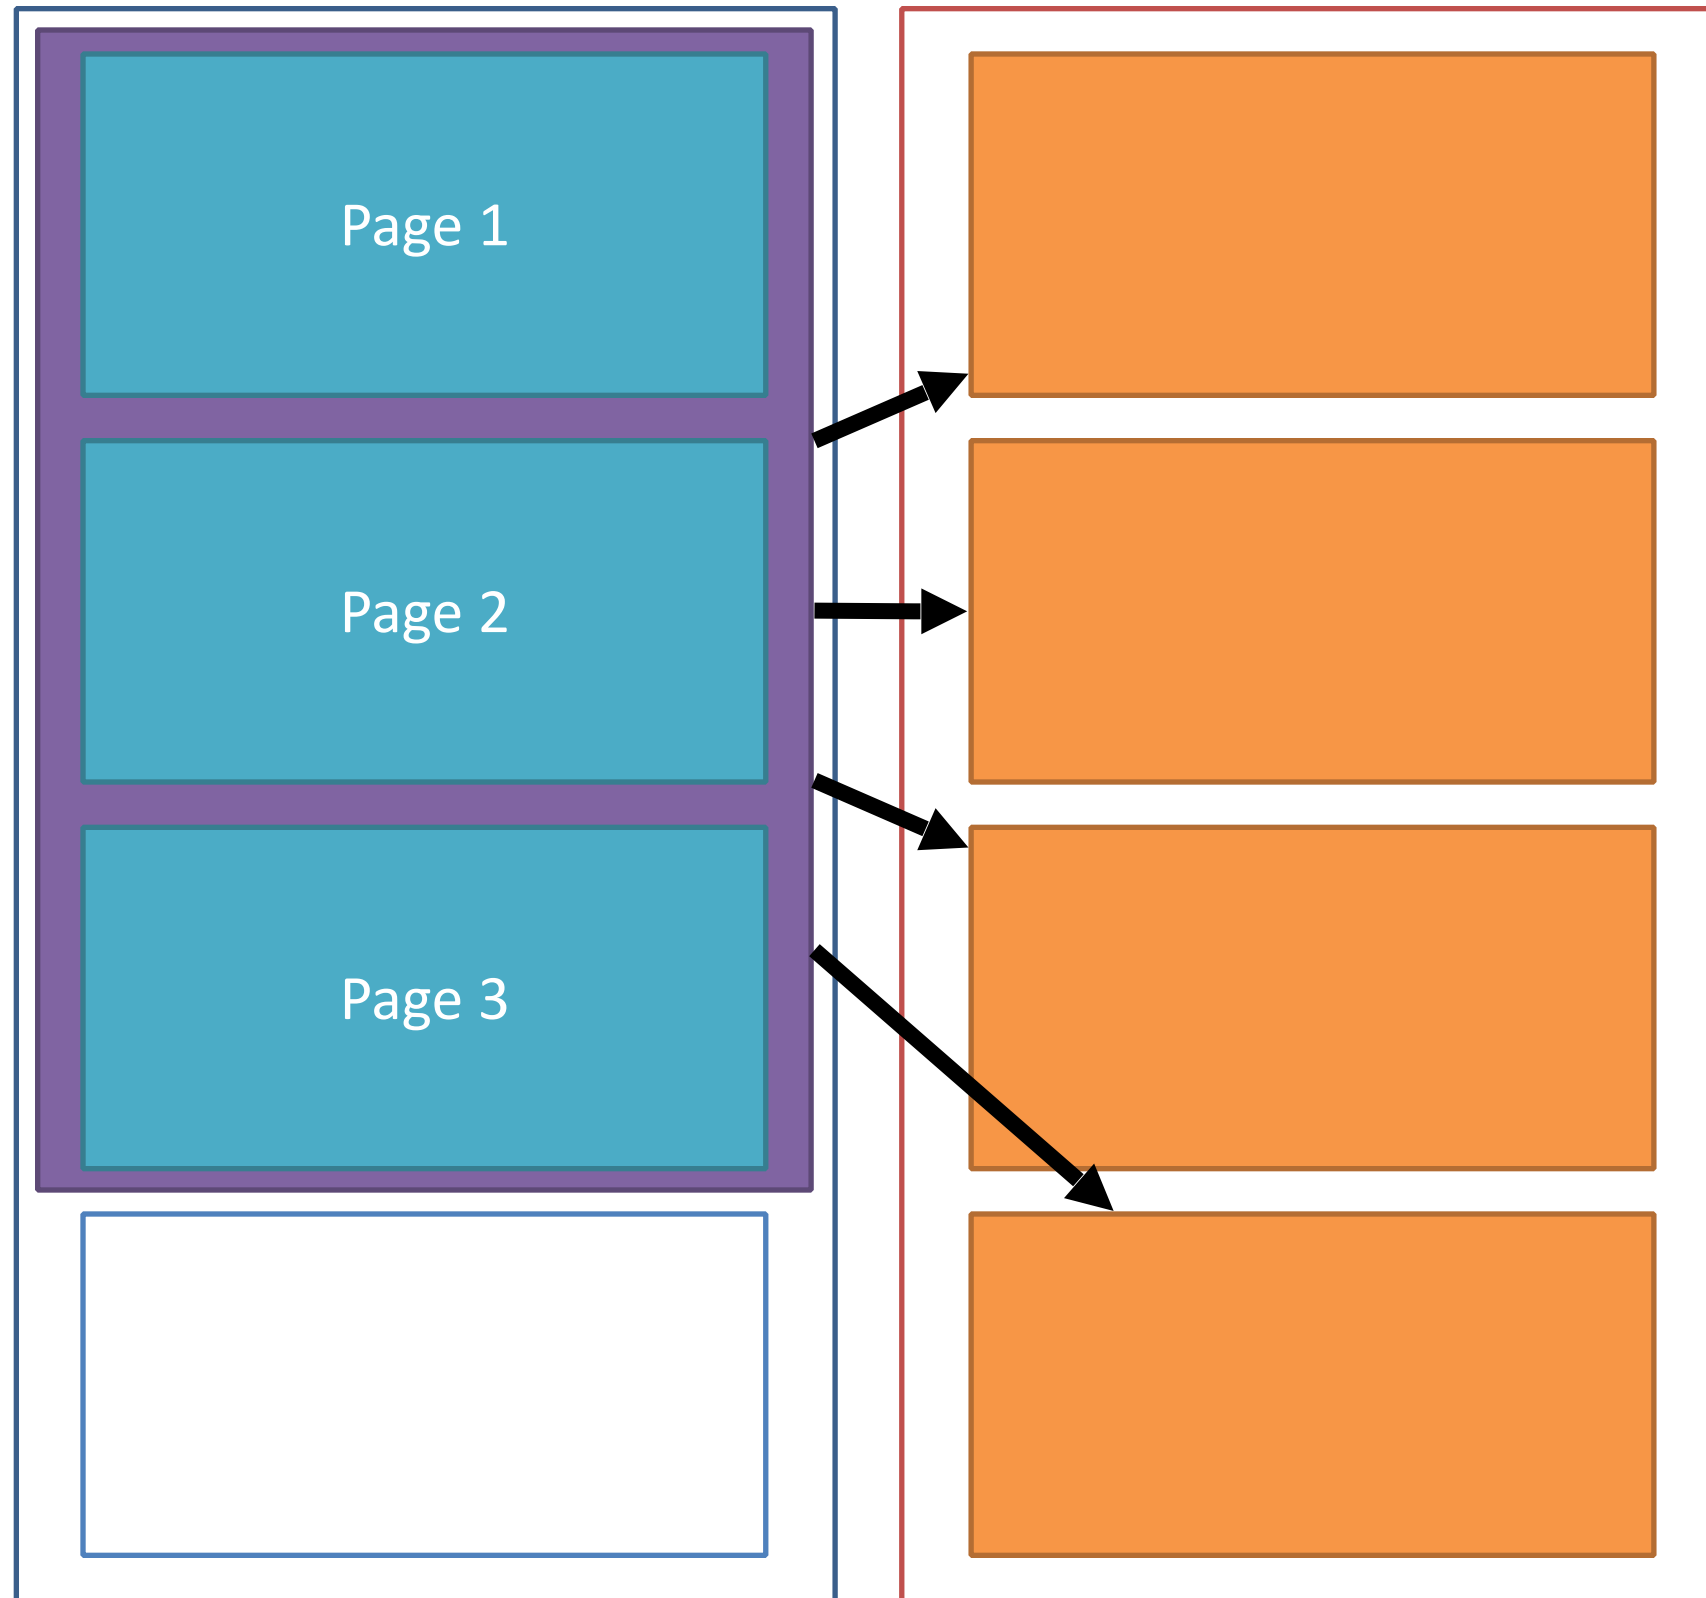
Steps:

1. Get **k** pages of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.

Steps:

1. Get **k** pages of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Chunk Nested Loops Join

Sailors

Reserves

Page 1

Page 2

Page 3

Page 4

Key idea:

Take **k pages** of S
and match with
each page of R.

Steps:

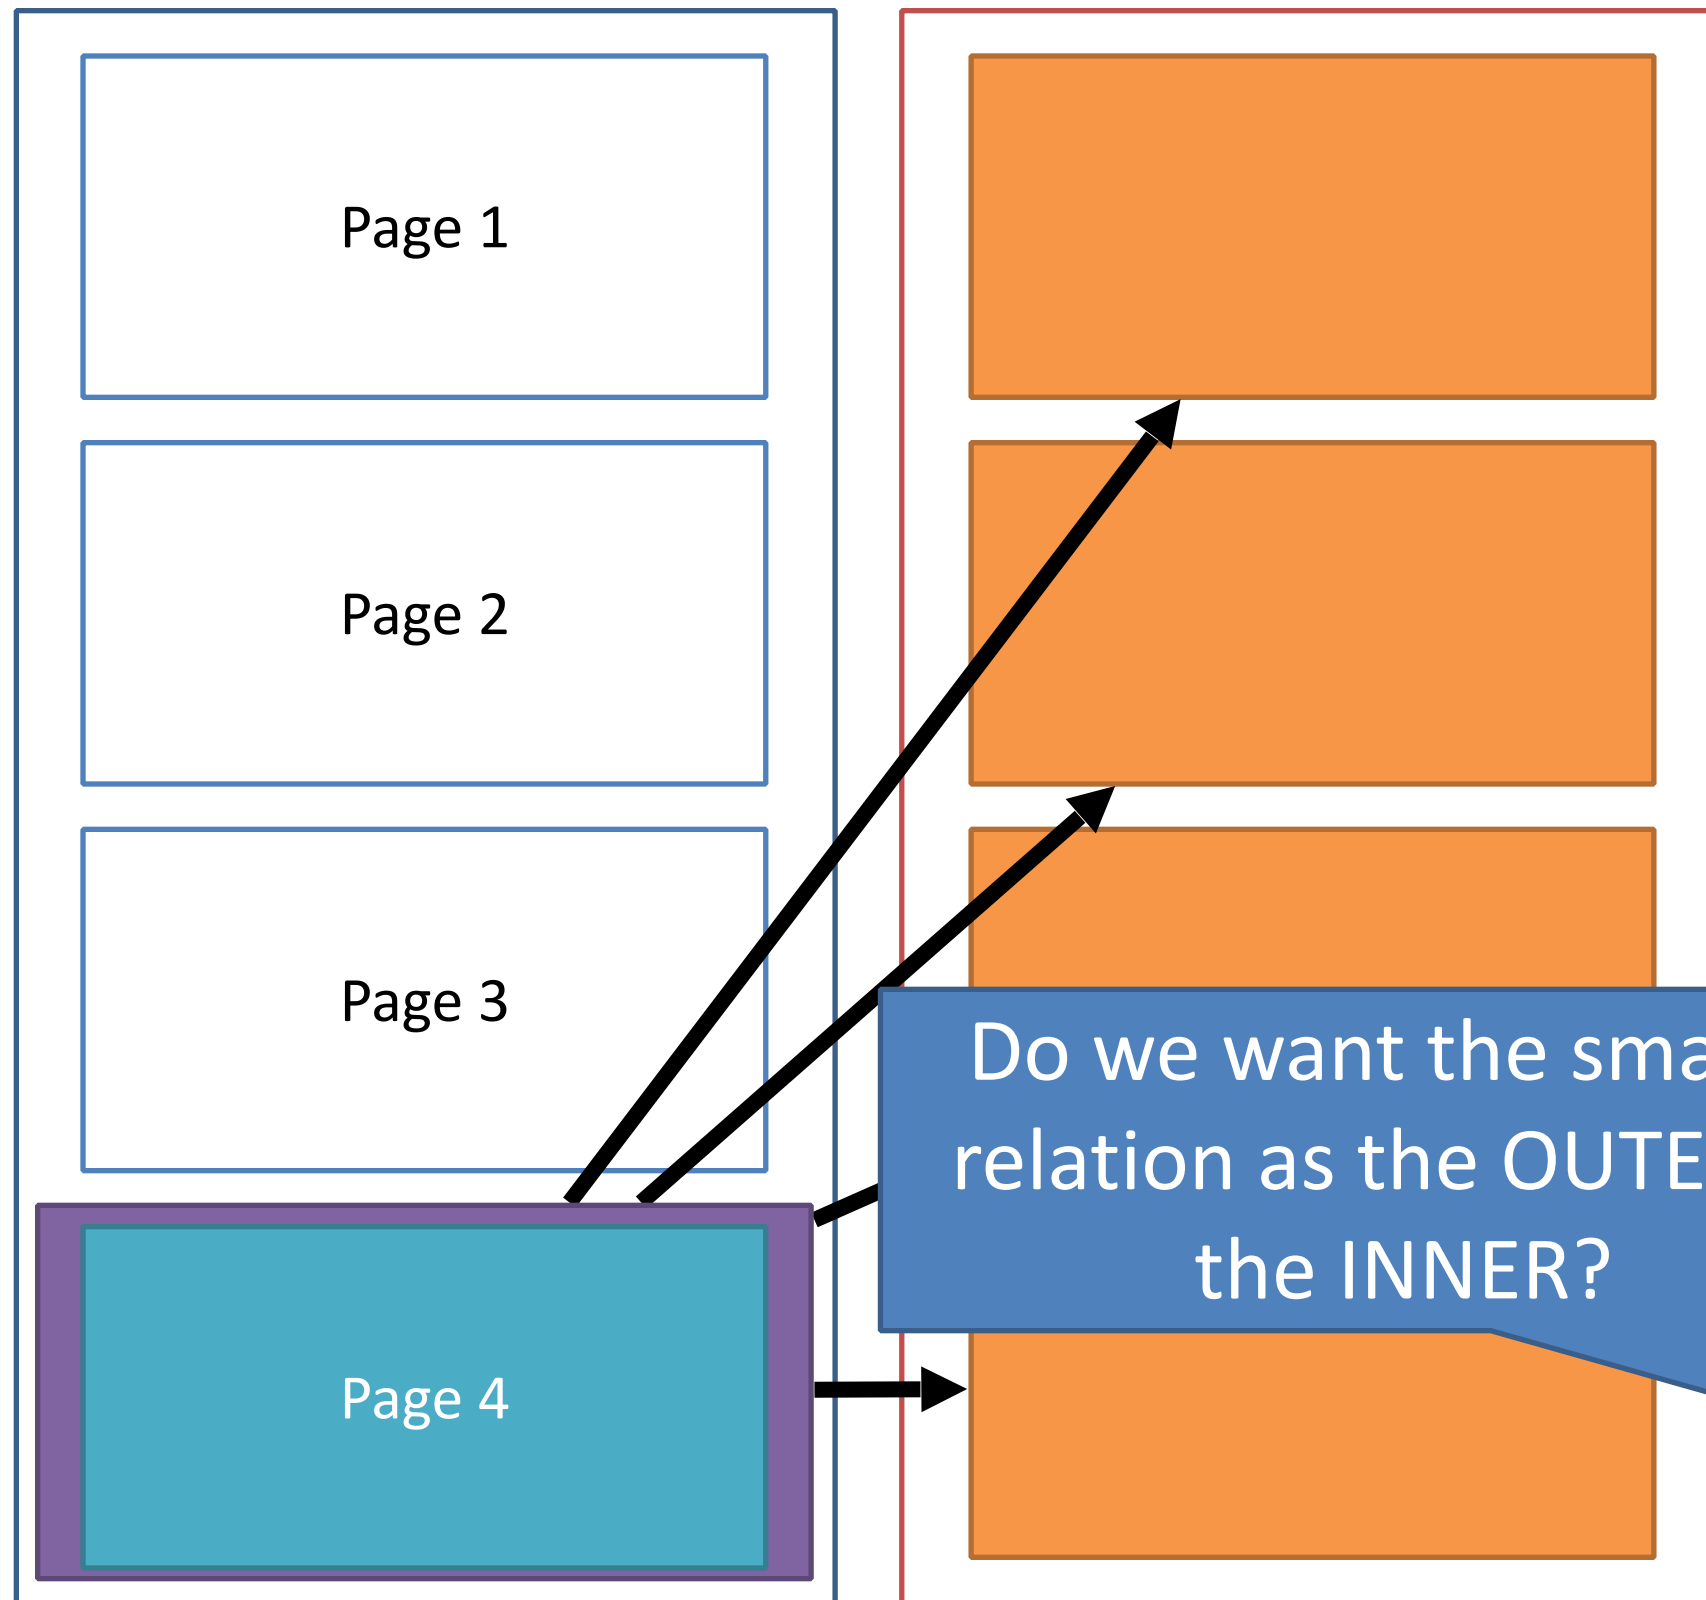
1. Get **k** pages of S.
2. Iterate through
each page in R.
3. Compare tuples in
each.

Notation: $[S]$ == “# pages in S” ; $|S|$ == “# tuples in S”

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.

Steps:

1. Get **k** pages of S.
2. Iterate through each page in R.

Compare tuples in each.

:

$$[S] + ([S] / k) * [R]$$

Do question #2 (a & b)

IMPORTANT

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R, then
merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

(name = Bob, sid = 1)

(name = Sam, sid = 3)

(name = Sue, sid = 7)

(name = Jill, sid = 2)

(name = Joe, sid = 12)

(name = Sue, sid = 8)

(name = Yue, sid = 4)

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

(name = Bob, sid = 1)

(name = Jill, sid = 2)

(name = Sam, sid = 3)

(name = Yue, sid = 4)

(name = Sue, sid = 7)

(name = Sue, sid = 8)

(name = Joe, sid = 12)

...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)

(name = Bob, sid = 1)

(name = Jill, sid = 2)

(name = Sam, sid = 3)

(name = Yue, sid = 4)

(name = Sue, sid = 7)

(name = Sue, sid = 8)

(name = Joe, sid = 12)

...

(sid = 1, bid = 4)

(sid = 1, bid = 7)

(sid = 3, bid = 6)

(sid = 4, bid = 3)

(sid = 8, bid = 1)

(sid = 8, bid = 13)

(sid = 8, bid = 15)

(sid = 12, bid = 1)

...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)



(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...
...
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...
...
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...



(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(name = Sue, sid = 8)
(name = Joe, sid = 12)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)

...

...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)

...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)

...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)

...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)

...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)

...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)

...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)

...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)

...

Notation: $[S]$ == “# pages in S ” ; $|S|$ == “# tuples in S ”

Sort-Merge Join

Sailors

Reserves

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

I/Os:

$$\sim 5([S] + [R])$$

Sorting: 4([S]+[R])

Merging: [S]+[R]

Optimizing Sort-Merge Join

Sailors

Reserves

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Key idea:

Internal Sort on
both. Perform
merge on all runs!

Steps:

1. Internal sort S and R. (Pass 0)
2. Merge all runs.

Optimizing Sort-Merge Join

Sailors

Reserves

Key idea:

Internal Sort on
both. Perform
merge on all runs!

Steps:

1. Internal sort S and R. (Pass 0)
2. Merge all runs.

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Yue, sid = 4)
(name = Sue, sid = 8)
(name = Jack, sid = 18)

(name = Cat, sid = 22)

...

(name = Sam, sid = 3)
(name = Sue, sid = 7)
(name = Joe, sid = 12)

...

(sid = 1, bid = 4)

(sid = 1, bid = 7)

(sid = 4, bid = 3)

(sid = 8, bid = 1)

(sid = 8, bid = 13)

(sid = 12, bid = 1)

...

(sid = 3, bid = 6)

(sid = 8, bid = 15)

...

Optimizing Sort-Merge Join

Sailors

Reserves

Key idea:

Internal Sort on both.
Perform merge on all runs!

Steps:

1. Internal sort S and R.
(Pass 0)

merge all runs.

NOTE: What does this
assume about the number
of runs?

$\sim 3([S] + [R])$

Pass 0: $2([S] + [R])$

Merging: $[S] + [R]$

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Yue, sid = 4)
(name = Sue, sid = 8)
(name = Jack, sid = 18)

(name = Cat, sid = 22)
...


(name = Sam, sid = 3)
(name = Sue, sid = 7)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 4, bid = 3)
(sid = 8, bid = 1)
(sid = 8, bid = 13)


(sid = 12, bid = 1)
...

Hash-Join

Sailors



Reserves



Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
2. Re-Hash, collect

Hash-Join

Reserves

Sailors

Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
2. Re-Hash, collect

(name = Bob, sid = 1)

(name = Sam, sid = 3)

(name = Sue, sid = 7)

(name = Jill, sid = 2)

(name = Joe, sid = 12)

(name = Sue, sid = 8)

(name = Yue, sid = 4)

Hash function: sid mod 4 Hash-Join

Sailors

Reserves

Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
2. Re-Hash, collect

(name = Joe, sid = 12)
(name = Sue, sid = 8)
(name = Yue, sid = 4)
...

(name = Bob, sid = 1)
...

(name = Jill, sid = 2)
...

(name = Sue, sid = 7)
(name = Sam, sid = 3)
...

(sid = 12, bid = 1)
(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 1, bid = 4)
(sid = 1, bid = 7)
...


(sid = 3, bid = 6)
...

...

Hash-Join

Sailors

Reserves



Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
2. Re-Hash, collect

Notation: $[S]$ == “# pages in S ” ; $|S|$ == “# tuples in S ”

Hash-Join

Sailors

Reserves

Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
using same hash fn, collect

Partition: $2([S] + [R])$

Re-Hash: $[S] + [R]$

NOTE: This is no different
from what we previously
assumed about hashing.

NOTE: What are we
assuming about the size of
partitions?

(name = Joe, sid = 12)
(name = Sue, sid = 8)
(name = Yue, sid = 4)
...

(name = Bob, ...)

(name = Jill, sid = 2)
...

(name = Sue, sid = 7)
(name = Sam, sid = 3)
...

(sid = 12, bid = 1)
(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

...

...

...


Hybrid Hashing (Hash Join Made Better)

- The hybrid hash join takes advantage of more available memory. During the partitioning phase, the hybrid hash join uses the available memory for two purposes:
 - To hold the current output buffer page for each of the k partitions
 - To hold an entire partition in-memory, known as "partition 0"
- Because partition 0 is never written to or read from disk, the hybrid hash join typically performs fewer I/O operations than regular hash join. Note that this algorithm is memory-sensitive, because there are two competing demands for memory (the hash table for partition 0, and the output buffers for the remaining partitions). Choosing too large a hash table might cause the algorithm to recurse because one of the non-zero partitions is too large to fit into memory.

Hybrid Hashing




Idea: keep one of the hash buckets in memory!



Cost savings: hybrid hashing



- Now:



Join Cheatsheet

Notation: $[S]$ == "# pages in S" ;

$|S|$ == "# tuples in S"

- Chunk nested loop join
 - Take k pages of S and match with each page of R .
 - Total Cost: $[S] + ([S] / k) * [R]$
- Sort merge join
 - Sort S and R on join column, then merge them!
 - Total Cost: $\sim 5([S] + [R])$
- Hash join
 - Partition S and R using same hash fn, then collect same partitions
 - Total Cost: $\sim 3([S] + [R])$
 - Assuming $\text{len}(\text{partition}) \leq B$ pages

Joins that we didn't go over!

- Index Nested Loop Join
 - More relevant when we start talking about B+ Trees, Indices, and Text Search
 - Look at lecture slides for clear pictures and understanding

When is a chunk-nested loops join the best?

When is a chunk-nested loops join the best?

- Not using an equality predicate
- Join is just a cross product

When is a sort-merge join the best?

When is a sort-merge join the best?

- Skewed input data
- Small memory size
- Want sorted output/have sorted input

When is a hash-join the best?

When is a hash-join the best?

- One partition large, the other small (can keep in memory) —> Hybrid Hashing

Do questions #1 and #2
(c & d)

IMPORTANT